

# Package ‘forecast’

June 4, 2011

**Title** Forecasting functions for time series

**Description** Methods and tools for displaying and analysing univariate time series forecasts including exponential smoothing via state space models and automatic ARIMA modelling.

**Version** 2.19

**Date** 2011-06-04

**Depends** R (>= 2.0.0), graphics, stats, tseries, fracdiff

**LazyData** yes

**LazyLoad** yes

**Author** Rob J Hyndman <Rob.Hyndman@monash.edu>

**Maintainer** Rob J Hyndman <Rob.Hyndman@monash.edu>

**License** GPL (>= 2)

**URL** <http://robjhyndman.com/software/forecast/>

**Repository** CRAN

**Date/Publication** 2011-06-04 08:07:23

## R topics documented:

accuracy . . . . .	2
Acf . . . . .	3
arfima . . . . .	4
Arima . . . . .	6
arima.errors . . . . .	8
auto.arima . . . . .	9
BoxCox . . . . .	10
croston . . . . .	11
CV . . . . .	13

decompose . . . . .	14
dm.test . . . . .	15
ets . . . . .	16
fitted.Arima . . . . .	18
forecast . . . . .	19
forecast.Arima . . . . .	20
forecast.ets . . . . .	22
forecast.HoltWinters . . . . .	23
forecast.lm . . . . .	24
forecast.stl . . . . .	26
forecast.StructTS . . . . .	27
gas . . . . .	29
gold . . . . .	29
logLik.ets . . . . .	30
ma . . . . .	31
meanf . . . . .	31
monthdays . . . . .	33
na.interp . . . . .	33
naive . . . . .	34
ndiffs . . . . .	35
plot.ets . . . . .	36
plot.forecast . . . . .	37
rwf . . . . .	38
seasadj . . . . .	40
seasonaldummy . . . . .	41
seasonplot . . . . .	42
ses . . . . .	43
simulate.ets . . . . .	45
sindexf . . . . .	46
splinef . . . . .	47
thetaf . . . . .	48
tsdisplay . . . . .	50
tslm . . . . .	51
wineind . . . . .	52
woolymq . . . . .	52

<b>Index</b>	<b>54</b>
--------------	-----------

---

accuracy

*Accuracy measures for forecast model*

---

## Description

Returns range of summary measures of the forecast accuracy. If  $x$  is provided, the function measures out-of-sample forecast accuracy based on  $x-f$ . If  $x$  is not provided, the function produces in-sample accuracy measures of the one-step forecasts based on  $f["x"]$ -fitted( $f$ ). All measures are defined and discussed in Hyndman and Koehler (2006).

**Usage**

```
accuracy(f, x, test=1:length(x), lambda=NULL)
```

**Arguments**

f	An object of class "forecast", or a numerical vector containing forecasts.
x	An optional numerical vector containing actual values of the same length as object.
test	Indicator of which elements of x and f to test.
lambda	Box-Cox parameter. If present, function backtransforms data and forecasts using <a href="#">InvBoxCox</a> .

**Value**

Vector giving forecast accuracy measures.

**Author(s)**

Rob J Hyndman

**References**

Hyndman, R.J. and Koehler, A.B. (2006) "Another look at measures of forecast accuracy". *International Journal of Forecasting*, **22**(4).

**Examples**

```
fit1 <- rwf(EuStockMarkets[1:200,1],h=100)
fit2 <- meanf(EuStockMarkets[1:200,1],h=100)
accuracy(fit1)
accuracy(fit2)
accuracy(fit1,EuStockMarkets[201:300,1])
accuracy(fit2,EuStockMarkets[201:300,1])
plot(fit1)
lines(EuStockMarkets[1:300,1])
```

---

Acf

---

(Partial) Autocorrelation Function Estimation

---

**Description**

Largely wrappers for the [acf](#) function in the stats package. The main difference is that Acf does not plot a spike at lag 0 (which is redundant). Pacf is included for consistency.

**Usage**

```
Acf(x, lag.max=NULL, type = c("correlation", "partial"), plot=TRUE, main=NULL, ...)
Pacf(x, main=NULL, ...)
```

**Arguments**

<code>x</code>	a univariate time series
<code>lag.max</code>	maximum lag at which to calculate the acf. Default is $10 \cdot \log_{10}(N/m)$ where $N$ is the number of observations and $m$ the number of series. Will be automatically limited to one less than the number of observations in the series.
<code>type</code>	character string giving the type of acf to be computed. Allowed values are "correlation" (the default) or "partial".
<code>plot</code>	logical. If TRUE (the default) the acf is plotted.
<code>main</code>	Title for plot
<code>...</code>	Additional arguments passed to <code>acf</code> .

**Details**

See the `acf` function in the stats package.

**Value**

See the `acf` function in the stats package.

**Author(s)**

Rob J Hyndman

**See Also**

`acf`

**Examples**

```
Acf(wineind)
Pacf(wineind)
```

---

arfima

*Fit a fractionally differenced ARFIMA model*

---

**Description**

An ARFIMA(p,d,q) model is selected and estimated automatically using the Hyndman-Khandakar (2008) algorithm to select p and q and the Haslett and Raftery (1989) algorithm to estimate the parameters including d.

**Usage**

```
arfima(x, drange = c(0, 0.5), estim = c("mle", "ls"), ...)
```

### Arguments

<code>x</code>	a univariate time series (numeric vector).
<code>drange</code>	Allowable values of <code>d</code> to be considered. Default of <code>c(0, 0.5)</code> ensures a stationary model is returned.
<code>estim</code>	If <code>estim=="ls"</code> , then the ARMA parameters are calculated using the Haslett-Raftery algorithm. If <code>estim=="mle"</code> , then the ARMA parameters are calculated using full MLE via the <code>arima</code> function.
<code>...</code>	Other arguments passed to <code>auto.arima</code> when selecting <code>p</code> and <code>q</code> .

### Details

This function combines `fracdiff` and `auto.arima` to automatically select and estimate an ARFIMA model. The fractional differencing parameter is chosen first assuming an ARFIMA(2,d,0) model. Then the data are fractionally differenced using the estimated `d` and an ARMA model is selected for the resulting time series using `auto.arima`. Finally, the full ARFIMA(p,d,q) model is re-estimated using `fracdiff`. If `estim=="mle"`, the ARMA coefficients are refined using `arima`.

### Value

A list object of S3 class `"fracdiff"`, which is described in the `fracdiff` documentation. A few additional objects are added to the list including `x` (the original time series), and the `residuals` and `fitted` values.

### Author(s)

Rob J Hyndman and Farah Yasmeen

### References

J. Haslett and A. E. Raftery (1989) Space-time Modelling with Long-memory Dependence: Assessing Ireland's Wind Power Resource (with discussion); *Applied Statistics* **38**, 1-50.

Hyndman, R.J. and Khandakar, Y. (2008) "Automatic time series forecasting: The forecast package for R", *Journal of Statistical Software*, **26**(3).

### See Also

`fracdiff`, `auto.arima`, `forecast.fracdiff`.

### Examples

```
x <- fracdiff.sim( 100, ma = -.4, d = .3)$series
fit <- arfima(x)
tsdisplay(residuals(fit))
```

**Description**

Largely a wrapper for the `arima` function in the `stats` package. The main difference is that this function allows a drift term. It is also possible to take an ARIMA model from a previous call to `Arima` and re-apply it to the data `x`.

**Usage**

```
Arima(x, order = c(0, 0, 0), seasonal = list(order = c(0, 0, 0), period = NA),
      xreg = NULL, include.mean = TRUE, include.drift = FALSE,
      transform.pars = TRUE, fixed = NULL, init = NULL,
      method = c("CSS-ML", "ML", "CSS"), n.cond,
      optim.control = list(), kappa = 1e6, model=NULL)
```

**Arguments**

<code>x</code>	a univariate time series
<code>order</code>	A specification of the non-seasonal part of the ARIMA model: the three components (p, d, q) are the AR order, the degree of differencing, and the MA order.
<code>seasonal</code>	A specification of the seasonal part of the ARIMA model, plus the period (which defaults to <code>frequency(x)</code> ). This should be a list with components <code>order</code> and <code>period</code> , but a specification of just a numeric vector of length 3 will be turned into a suitable list with the specification as the <code>order</code> .
<code>xreg</code>	Optionally, a vector or matrix of external regressors, which must have the same number of rows as <code>x</code> .
<code>include.mean</code>	Should the ARIMA model include a mean term? The default is <code>TRUE</code> for undifferenced series, <code>FALSE</code> for differenced ones (where a mean would not affect the fit nor predictions).
<code>include.drift</code>	Should the ARIMA model include a linear drift term? (i.e., a linear regression with ARIMA errors is fitted.) The default is <code>FALSE</code> .
<code>transform.pars</code>	Logical. If true, the AR parameters are transformed to ensure that they remain in the region of stationarity. Not used for <code>method = "CSS"</code> .
<code>fixed</code>	optional numeric vector of the same length as the total number of parameters. If supplied, only NA entries in <code>fixed</code> will be varied. <code>transform.pars = TRUE</code> will be overridden (with a warning) if any AR parameters are fixed. It may be wise to set <code>transform.pars = FALSE</code> when fixing MA parameters, especially near non-invertibility.
<code>init</code>	optional numeric vector of initial parameter values. Missing values will be filled in, by zeroes except for regression coefficients. Values already specified in <code>fixed</code> will be ignored.

<code>method</code>	Fitting method: maximum likelihood or minimize conditional sum-of-squares. The default (unless there are missing values) is to use conditional-sum-of-squares to find starting values, then maximum likelihood.
<code>n.cond</code>	Only used if fitting by conditional-sum-of-squares: the number of initial observations to ignore. It will be ignored if less than the maximum lag of an AR term.
<code>optim.control</code>	List of control parameters for <code>optim</code> .
<code>kappa</code>	the prior variance (as a multiple of the innovations variance) for the past observations in a differenced model. Do not reduce this.
<code>model</code>	Output from a previous call to <code>Arima</code> . If <code>model</code> is passed, this same model is fitted to <code>x</code> without re-estimating any parameters.

### Details

See the [arima](#) function in the stats package.

### Value

See the [arima](#) function in the stats package. The additional objects returned are

<code>x</code>	The time series data
<code>xreg</code>	The regressors used in fitting (when relevant).

### Author(s)

Rob J Hyndman

### See Also

[arima](#)

### Examples

```
fit <- Arima(WWWusage, order=c(3,1,0))
plot(forecast(fit, h=20))

air.model <- Arima(window(AirPassengers, end=1956+11/12), order=c(0,1,1), seasonal=list(order=c(1,1,0,0)))
plot(forecast(air.model, h=48))
lines(AirPassengers)

air.model2 <- Arima(window(AirPassengers, start=1957), model=air.model)
outofsample <- fitted(air.model2)
# in-sample one-step forecasts
accuracy(air.model)
# out-of-sample one-step forecasts
accuracy(air.model2)
# out-of-sample multi-step forecasts
accuracy(forecast(air.model, h=48), outofsample)
```

---

arima.errors	<i>ARIMA errors</i>
--------------	---------------------

---

## Description

Returns original time series after adjusting for regression variables. These are not the same as the residuals. If there are no regression variables in the ARIMA model, then the errors will be identical to the original series. If there are regression variables in the ARIMA model, then the errors will be equal to the original series minus the effect of the regression variables, but leaving in the serial correlation that is modelled with the AR and MA terms. If you want the "residuals", then use `residuals(z)` ..

## Usage

```
arima.errors(z)
```

## Arguments

<code>z</code>	Fitted ARIMA model from <code>arima</code>
----------------	--

## Value

A time series containing the "errors".

## Author(s)

Rob J Hyndman

## See Also

`arima`, `residuals`

## Examples

```
ukdeaths.fit <- Arima(UKDriverDeaths,c(1,0,1),c(0,1,1),xreg=Seatbelts[, "law"])
ukdeaths.errors <- arima.errors(ukdeaths.fit)
par(mfrow=c(2,1))
plot(UKDriverDeaths)
plot(ukdeaths.errors)
```



auto.arima

*Fit best ARIMA model to univariate time series***Description**

Returns best ARIMA model according to either AIC, AICc or BIC value. The function conducts a search over possible model within the order constraints provided.

**Usage**

```
auto.arima(x, d = NA, D = NA, max.p = 5, max.q = 5,
           max.P = 2, max.Q = 2, max.order = 5,
           start.p=2, start.q=2, start.P=1, start.Q=1,
           stationary = FALSE, ic = c("aic", "aicc", "bic"),
           stepwise=TRUE, trace=FALSE,
           approximation=length(x)>100 | frequency(x)>12, xreg=NULL,
           test=c("kpss", "adf", "pp"), allowdrift=TRUE)
```

**Arguments**

x	a univariate time series
d	Order of first-differencing. If missing, will choose a value based on KPSS test.
D	Order of seasonal-differencing. If missing, will choose a value based on CH test.
max.p	Maximum value of p
max.q	Maximum value of q
max.P	Maximum value of P
max.Q	Maximum value of Q
max.order	Maximum value of p+q+P+Q if model selection is not stepwise.
start.p	Starting value of p in stepwise procedure.
start.q	Starting value of q in stepwise procedure.
start.P	Starting value of P in stepwise procedure.
start.Q	Starting value of Q in stepwise procedure.
stationary	If TRUE, restricts search to stationary models.
ic	Information criterion to be used in model selection.
stepwise	If TRUE, will do stepwise selection (faster). Otherwise, it searches over all models. Non-stepwise selection can be very slow, especially for seasonal models.
trace	If TRUE, the list of ARIMA models considered will be reported.
approximation	If TRUE, estimation is via conditional sums of squares and the information criteria used for model selection are approximated. The final model is still computed using maximum likelihood estimation. Approximation should be used for long time series or a high seasonal period to avoid excessive computation times.

<code>xreg</code>	Optionally, a vector or matrix of external regressors, which must have the same number of rows as <code>x</code> .
<code>test</code>	Type of unit root test to use. See <a href="#">ndiffs</a> for details.
<code>allowdrift</code>	If TRUE, models with drift terms are considered.

### Details

Non-stepwise selection can be slow, especially for seasonal data. Non-seasonal differences chosen using the KPSS test. Seasonal differences chosen using a variation on the Canova-Hansen test. Stepwise algorithm outlined in Hyndman and Khandakar (2008).

### Value

Same as for [arima](#)

### Author(s)

Rob J Hyndman

### References

Hyndman, R.J. and Khandakar, Y. (2008) "Automatic time series forecasting: The forecast package for R", *Journal of Statistical Software*, **26**(3).

### See Also

[Arima](#)

### Examples

```
fit <- auto.arima(WWWusage)
plot(forecast(fit, h=20))
```

---

BoxCox

*Box Cox Transformation*

---

### Description

BoxCox() returns a transformation of the input variable using a Box-Cox transformation. InvBoxCox() reverses the transformation.

### Usage

```
BoxCox(x, lambda)
InvBoxCox(x, lambda)
```

**Arguments**

x	a numeric vector or time series
lambda	transformation parameter

**Details**

The Box-Cox transformation is given by

$$f_{\lambda}(x) = \frac{x^{\lambda} - 1}{\lambda}$$

if  $\lambda \neq 0$ . For  $\lambda = 0$ ,

$$f_0(x) = \log(x)$$

.

**Value**

a numeric vector of the same length as x.

**Author(s)**

Rob J Hyndman

**References**

Box, G. E. P. and Cox, D. R. (1964) An analysis of transformations. *JRSS B* **26** 211–246.

**Examples**

```
lynx.sqrt <- BoxCox(lynx, 0.5)
lynx.fit <- ar(lynx.sqrt)
plot(forecast(lynx.fit, h=20), lambda=0.5)
```

---

croston

*Forecasts for intermittent demand using Croston's method*


---

**Description**

Returns forecasts and other information for Croston's forecasts applied to x.

**Usage**

```
croston(x, h=10, alpha=0.1)
```

**Arguments**

x	a numeric vector or time series
h	Number of periods for forecasting.
alpha	Value of alpha. Default value is 0.1.

## Details

Based on Croston's (1972) method for intermittent demand forecasting, also described in Shenstone and Hyndman (2005). Croston's method involves using simple exponential smoothing (SES) on the non-zero elements of the time series and a separate application of SES to the times between non-zero elements of the time series. The smoothing parameters of the two applications of SES are assumed to be equal and are denoted by `alpha`.

Note that prediction intervals are not computed as Croston's method has no underlying stochastic model.

## Value

An object of class `"forecast"` is a list containing at least the following elements:

<code>model</code>	A list containing information about the fitted model. The first element gives the SES model used for non-zero demands. The second element gives the SES model used for times between non-zero demands. Both models are of class <code>forecast</code> .
<code>method</code>	The name of the forecasting method as a character string
<code>mean</code>	Point forecasts as a time series
<code>x</code>	The original time series (either <code>object</code> itself or the time series used to create the model stored as <code>object</code> ).
<code>residuals</code>	Residuals from the fitted model. That is <code>x</code> minus fitted values.
<code>fitted</code>	Fitted values (one-step forecasts)

The function `summary` is used to obtain and print a summary of the results, while the function `plot` produces a plot of the forecasts.

The generic accessor functions `fitted.values` and `residuals` extract useful features of the value returned by `croston` and associated functions.

## Author(s)

Rob J Hyndman

## References

- Croston, J. (1972) "Forecasting and stock control for intermittent demands", *Operational Research Quarterly*, **23**(3), 289-303.
- Shenstone, L., and Hyndman, R.J. (2005) "Stochastic models underlying Croston's method for intermittent demand forecasting". *Journal of Forecasting*, **24**, 389-402.

## See Also

[ses](#).

**Examples**

```
x <- rpois(20, lambda=.3)
fcast <- croston(x)
plot(fcast)
```

---

CV

*Cross-validation statistic*

---

**Description**

Computes cross-validation statistic, AIC, corrected AIC, BIC and adjusted  $R^2$  values for a linear model.

**Usage**

```
CV(obj)
```

**Arguments**

`obj`                      output from `lm` or `tslm`

**Value**

Numerical vector containing CV, AIC, AICc, BIC and AdjR2 values.

**Author(s)**

Rob J Hyndman

**See Also**

[AIC](#)

**Examples**

```
y <- ts(rnorm(120, 0, 3) + 20*sin(2*pi*(1:120)/12), frequency=12)
fit1 <- tslm(y ~ trend + season)
fit2 <- tslm(y ~ season)
CV(fit1)
CV(fit2)
```

decompose

*Classical Seasonal Decomposition by Moving Averages***Description**

Decompose a time series into seasonal, trend and irregular components using moving averages. Deals with additive or multiplicative seasonal component.

**Usage**

```
decompose(x, type = c("additive", "multiplicative"), filter = NULL)
```

**Arguments**

<code>x</code>	A time series.
<code>type</code>	The type of seasonal component. Can be abbreviated.
<code>filter</code>	A vector of filter coefficients in reverse time order (as for AR or MA coefficients), used for filtering out the seasonal component. If <code>NULL</code> , a moving average with symmetric window is performed.

**Details**

The additive model used is:

$$Y_t = T_t + S_t + e_t$$

The multiplicative model used is:

$$Y_t = T_t S_t e_t$$

The function first determines the trend component using a moving average (if `filter` is `NULL`, a symmetric window with equal weights is used), and removes it from the time series. Then, the seasonal figure is computed by averaging, for each time unit, over all periods. The seasonal figure is then centered. Finally, the error component is determined by removing trend and seasonal figure (recycled as needed) from the original time series.

**Value**

An object of class `"decomposed.ts"` with following components:

<code>seasonal</code>	The seasonal component (i.e., the repeated seasonal figure)
<code>figure</code>	The estimated seasonal figure only
<code>trend</code>	The trend component
<code>random</code>	The remainder part
<code>type</code>	The value of <code>type</code>

**Note**

This function is identical to the `decompose` function in the `stats` package except that the seasonal component is not incorrectly truncated.

**Author(s)**

David Meyer <David.Meyer@wu-wien.ac.at>. Revised by Rob J Hyndman <Rob.Hyndman@monash.edu>.

**References**

M. Kendall and A. Stuart (1983) The Advanced Theory of Statistics, Vol.3, *Griffin*, 410–414.

**See Also**

[decompose](#)

**Examples**

```
m <- decompose(co2)
plot(m)
```

---

dm.test

---

*Diebold-Mariano test for predictive accuracy*


---

**Description**

The Diebold-Mariano test compares the forecast accuracy of two forecast methods. The null hypothesis is that they have the same forecast accuracy.

**Usage**

```
dm.test(e1, e2, alternative = c("two.sided", "less", "greater"), h = 1, power = 2)
```

**Arguments**

e1	Forecast errors from method 1.
e2	Forecast errors from method 2.
alternative	a character string specifying the alternative hypothesis, must be one of "two.sided" (default), "greater" or "less". You can specify just the initial letter.
h	The forecast horizon used in calculating e1 and e2.
power	The power used in the loss function. Usually 1 or 2.

**Value**

A list with class "htest" containing the following components:

statistic	the value of the DM-statistic.
parameter	the forecast horizon and loss function power used in the test.
alternative	a character string describing the alternative hypothesis.
p.value	the p-value for the test.
method	a character string with the value "Diebold-Mariano Test".
data.name	a character vector giving the names of the two error series.

**Author(s)**

George Athanasopoulos and Rob Hyndman

**References**

Diebold, F.X. and Mariano, R.S. (1995) Comparing predictive accuracy. *Journal of Business and Economic Statistics*, **13**, 253-263.

**Examples**

```
# Test on in-sample one-step forecasts
f1 <- ets(WWWusage)
f2 <- auto.arima(WWWusage)
accuracy(f1)
accuracy(f2)
dm.test(residuals(f1), residuals(f2), h=1)

# Test on out-of-sample one-step forecasts
f1 <- ets(WWWusage[1:80])
f2 <- auto.arima(WWWusage[1:80])
f1.out <- ets(WWWusage[81:100], model=f1)
f2.out <- Arima(WWWusage[81:100], model=f2)
accuracy(f1.out)
accuracy(f2.out)
dm.test(residuals(f1.out), residuals(f2.out), h=1)
```

---

ets

*Exponential smoothing state space model*

---

**Description**

Returns ets model applied to  $y$ .

**Usage**

```
ets(y, model="ZZZ", damped=NULL, alpha=NULL, beta=NULL, gamma=NULL, phi=NULL,
    additive.only=FALSE, lower=c(rep(0.0001, 3), 0.8), upper=c(rep(0.9999, 3), 0.98),
    opt.crit=c("lik", "amse", "mse", "sigma"), nmse=3,
    bounds=c("both", "usual", "admissible"), ic=c("aic", "aicc", "bic"),
    restrict=TRUE)
```

**Arguments**

<code>y</code>	a numeric vector or time series
<code>model</code>	Usually a three-character string identifying method using the framework terminology of Hyndman et al. (2002) and Hyndman et al. (2008). The first letter denotes the error type ("A", "M" or "Z"); the second letter denotes the trend type ("N", "A", "M" or "Z"); and the third letter denotes the season type ("N", "A", "M"



or "Z"). In all cases, "N"=none, "A"=additive, "M"=multiplicative and "Z"=automatically selected. So, for example, "ANN" is simple exponential smoothing with additive errors, "MAM" is multiplicative Holt-Winters' method with multiplicative errors, and so on. It is also possible for the model to be equal to the output from a previous call to `ets`. In this case, the same model is fitted to  $y$  without re-estimating any parameters.

<code>damped</code>	If TRUE, use a damped trend (either additive or multiplicative). If NULL, both damped and non-damped trends will be tried and the best model (according to the information criterion <code>ic</code> ) returned.
<code>alpha</code>	Value of alpha. If NULL, it is estimated.
<code>beta</code>	Value of beta. If NULL, it is estimated.
<code>gamma</code>	Value of gamma. If NULL, it is estimated.
<code>phi</code>	Value of phi. If NULL, it is estimated.
<code>additive.only</code>	If TRUE, will only consider additive models. Default is FALSE.
<code>lower</code>	Lower bounds for the parameters (alpha, beta, gamma, phi)
<code>upper</code>	Upper bounds for the parameters (alpha, beta, gamma, phi)
<code>opt.crit</code>	Optimization criterion. One of "mse" (Mean Square Error), "amse" (Average MSE over first <code>nmse</code> forecast horizons), "sigma" (Standard deviation of residuals), or "lik" (Log-likelihood, the default).
<code>nmse</code>	Number of steps for average multistep MSE ( $1 \leq \text{nmse} \leq 10$ ).
<code>bounds</code>	Type of parameter space to impose: "usual" indicates all parameters must lie between specified lower and upper bounds; "admissible" indicates parameters must lie in the admissible space; "both" (default) takes the intersection of these regions.
<code>ic</code>	Information criterion to be used in model selection.
<code>restrict</code>	If TRUE, the models with infinite variance will not be allowed.

## Details

Based on the classification of methods as described in Hyndman et al (2008).

The methodology is fully automatic. The only required argument for `ets` is the time series. The model is chosen automatically if not specified. This methodology performed extremely well on the M3-competition data. (See Hyndman, et al, 2002, below.)

## Value

An object of class "ets".

The generic accessor functions `fitted.values` and `residuals` extract useful features of the value returned by `ets` and associated functions.

## Author(s)

Rob J Hyndman

## References

- Hyndman, R.J., Koehler, A.B., Snyder, R.D., and Grose, S. (2002) "A state space framework for automatic forecasting using exponential smoothing methods", *International J. Forecasting*, **18**(3), 439–454.
- Hyndman, R.J., Akram, Md., and Archibald, B. (2008) "The admissible parameter space for exponential smoothing models". *Annals of Statistical Mathematics*, **60**(2), 407–426.
- Hyndman, R.J., Koehler, A.B., Ord, J.K., and Snyder, R.D. (2008) *Forecasting with exponential smoothing: the state space approach*, Springer-Verlag. <http://robjhyndman.com/expsmooth>.

## See Also

[HoltWinters](#), [rwf](#), [arima](#).

## Examples

```
fit <- ets(USAccDeaths)
plot(forecast(fit))
```

---

fitted.Arima

*One-step in-sample forecasts using ARIMA models*

---

## Description

Returns one-step forecasts for the data used in fitting the ARIMA model.

## Usage

```
fitted.Arima(object, ...)
```

## Arguments

<code>object</code>	An object of class "Arima". Usually the result of a call to <a href="#">arima</a> .
<code>...</code>	Other arguments.

## Value

An time series of the one-step forecasts.

## Author(s)

Rob J Hyndman

## See Also

[forecast.Arima](#).

## Examples

```
fit <- Arima(WWWusage,c(3,1,0))
plot(WWWusage)
lines(fitted(fit),col=2)
```

---

forecast

Forecasting time series

---

## Description

`forecast` is a generic function for forecasting from time series or time series models. The function invokes particular *methods* which depend on the class of the first argument.

For example, the function `forecast.Arima` makes forecasts based on the results produced by `arima`.

The function `forecast.ts` makes forecasts using exponential smoothing state space models.

## Usage

```
forecast(object,...)
## S3 method for class 'ts'
forecast(object, h, level=c(80,95), fan=FALSE, ...)
```

## Arguments

<code>object</code>	a time series or time series model for which forecasts are required
<code>h</code>	Number of periods for forecasting
<code>level</code>	Confidence level for prediction intervals.
<code>fan</code>	If TRUE, level is set to seq(50,99,by=1). This is suitable for fan plots.
<code>...</code>	Additional arguments affecting the forecasts produced. <code>forecast.ts</code> passes these to <code>forecast.ets</code>

## Details

The default behaviour is to use a model estimated using `ets`.

## Value

An object of class "forecast".

The function `summary` is used to obtain and print a summary of the results, while the function `plot` produces a plot of the forecasts and prediction intervals.

The generic accessor functions `fitted.values` and `residuals` extract various useful features of the value returned by `forecast$model`.

An object of class "forecast" is a list containing at least the following elements:

<code>model</code>	A list containing information about the fitted model
--------------------	--

method	The name of the forecasting method as a character string
mean	Point forecasts as a time series
lower	Lower limits for prediction intervals
upper	Upper limits for prediction intervals
level	The confidence values associated with the prediction intervals
x	The original time series (either <code>object</code> itself or the time series used to create the model stored as <code>object</code> ).
residuals	Residuals from the fitted model. That is <code>x</code> minus fitted values.
fitted	Fitted values (one-step forecasts)

**Author(s)**

Rob J Hyndman

**See Also**

Other functions which return objects of class "forecast" are [forecast.ets](#), [forecast.Arima](#), [forecast.HoltWinters](#), [forecast.StructTS](#), [meanf](#), [rwf](#), [splinef](#), [thetaf](#), [croston](#), [ses](#), [holt](#), [hw](#).

---

forecast.Arima

*Forecasting using ARIMA or ARFIMA models*


---

**Description**

Returns forecasts and other information for univariate ARIMA models.

**Usage**

```
## S3 method for class 'Arima'
forecast(object, h=ifelse(object$arima[5]>1, 2*object$arima[5], 10),
         level=c(80,95), fan=FALSE, xreg=NULL, ...)
## S3 method for class 'ar'
forecast(object, h=10, level=c(80,95), fan=FALSE, ...)
## S3 method for class 'fracdiff'
forecast(object, h=10, level=c(80,95), fan=FALSE, ...)
```

**Arguments**

object	An object of class "Arima", "ar" or "fracdiff". Usually the result of a call to <a href="#">arima</a> , <a href="#">auto.arima</a> , <a href="#">ar</a> , <a href="#">arfima</a> or <a href="#">fracdiff</a> .
h	Number of periods for forecasting. If <code>xreg</code> is used, <code>h</code> is ignored and the number of forecast periods is set to the number of rows of <code>xreg</code> .
level	Confidence level for prediction intervals.
fan	If TRUE, level is set to <code>seq(50, 99, by=1)</code> . This is suitable for fan plots.
xreg	Future values of an regression variables (for class <code>Arima</code> objects only).
...	Other arguments.

## Details

For Arima or ar objects, the function calls `predict.Arima` or `predict.ar` and constructs an object of class "forecast" from the results. For fracdiff objects, the calculations are all done within `forecast.fracdiff` using the equations given by Peiris and Perera (1988).

## Value

An object of class "forecast".

The function `summary` is used to obtain and print a summary of the results, while the function `plot` produces a plot of the forecasts and prediction intervals.

The generic accessor functions `fitted.values` and `residuals` extract useful features of the value returned by `forecast.Arima`.

An object of class "forecast" is a list containing at least the following elements:

<code>model</code>	A list containing information about the fitted model
<code>method</code>	The name of the forecasting method as a character string
<code>mean</code>	Point forecasts as a time series
<code>lower</code>	Lower limits for prediction intervals
<code>upper</code>	Upper limits for prediction intervals
<code>level</code>	The confidence values associated with the prediction intervals
<code>x</code>	The original time series (either <code>object</code> itself or the time series used to create the model stored as <code>object</code> ).
<code>residuals</code>	Residuals from the fitted model. That is <code>x</code> minus fitted values.
<code>fitted</code>	Fitted values (one-step forecasts)

## Author(s)

Rob J Hyndman

## References

Peiris, M. & Perera, B. (1988), On prediction with fractionally differenced ARIMA models, *Journal of Time Series Analysis*, **9**(3), 215-220.

## See Also

`predict.Arima`, `predict.ar`, `auto.arima`, `Arima`, `arima`, `ar`, `arfima`.

## Examples

```
fit <- Arima(WWWusage, c(3, 1, 0))
plot(forecast(fit))

x <- fracdiff.sim( 100, ma = -.4, d = .3)$series
fit <- arfima(x)
plot(forecast(fit, h=30))
```

forecast.ets

Forecasting using ETS models

**Description**

Returns forecasts and other information for univariate ETS models.

**Usage**

```
## S3 method for class 'ets'
forecast(object, h=ifelse(object$m>1, 2*object$m, 10), level=c(80,95),
fan=FALSE, simulate=FALSE, bootstrap=FALSE, npaths=5000,...)
```

**Arguments**

object	An object of class "ets". Usually the result of a call to <a href="#">ets</a> .
h	Number of periods for forecasting
level	Confidence level for prediction intervals.
fan	If TRUE, level is set to seq(50,99,by=1). This is suitable for fan plots.
simulate	If TRUE, prediction intervals produced by simulation rather than using analytic formulae.
bootstrap	If TRUE, and if simulate=TRUE, then simulation uses resampled errors rather than normally distributed errors.
npaths	Number of sample paths used in computing simulated prediction intervals.
...	Other arguments.

**Value**

An object of class "forecast".

The function `summary` is used to obtain and print a summary of the results, while the function `plot` produces a plot of the forecasts and prediction intervals.

The generic accessor functions `fitted.values` and `residuals` extract useful features of the value returned by `forecast.ets`.

An object of class "forecast" is a list containing at least the following elements:

model	A list containing information about the fitted model
method	The name of the forecasting method as a character string
mean	Point forecasts as a time series
lower	Lower limits for prediction intervals
upper	Upper limits for prediction intervals
level	The confidence values associated with the prediction intervals
x	The original time series (either <code>object</code> itself or the time series used to create the model stored as <code>object</code> ).
residuals	Residuals from the fitted model. That is <code>x</code> minus fitted values.
fitted	Fitted values (one-step forecasts)

**Author(s)**

Rob J Hyndman

**See Also**

[ets](#), [ses](#), [holt](#), [hw](#).

**Examples**

```
fit <- ets(USAccDeaths)
plot(forecast(fit, h=48))
```

---

```
forecast.HoltWinters
```

*Forecasting using Holt-Winters objects*

---

**Description**

Returns forecasts and other information for univariate Holt-Winters time series models.

**Usage**

```
## S3 method for class 'HoltWinters'
forecast(object, h=ifelse(frequency(object$x)>1, 2*frequency(object$x), 10),
  level=c(80, 95), fan=FALSE, ...)
```

**Arguments**

object	An object of class "HoltWinters". Usually the result of a call to <a href="#">HoltWinters</a> .
h	Number of periods for forecasting
level	Confidence level for prediction intervals.
fan	If TRUE, level is set to seq(50,99,by=1). This is suitable for fan plots.
...	Other arguments.

**Details**

This function calls [predict.HoltWinters](#) and constructs an object of class "forecast" from the results.

It is included for completeness, but the [ets](#) is recommended for use instead of [HoltWinters](#).

**Value**

An object of class "forecast".

The function `summary` is used to obtain and print a summary of the results, while the function `plot` produces a plot of the forecasts and prediction intervals.

The generic accessor functions `fitted.values` and `residuals` extract useful features of the value returned by `forecast.HoltWinters`.

An object of class "forecast" is a list containing at least the following elements:

<code>model</code>	A list containing information about the fitted model
<code>method</code>	The name of the forecasting method as a character string
<code>mean</code>	Point forecasts as a time series
<code>lower</code>	Lower limits for prediction intervals
<code>upper</code>	Upper limits for prediction intervals
<code>level</code>	The confidence values associated with the prediction intervals
<code>x</code>	The original time series (either <code>object</code> itself or the time series used to create the model stored as <code>object</code> ).
<code>residuals</code>	Residuals from the fitted model. That is <code>x</code> minus fitted values.
<code>fitted</code>	Fitted values (one-step forecasts)

**Author(s)**

Rob J Hyndman

**See Also**

[`predict.HoltWinters`](#), [`HoltWinters`](#).

**Examples**

```
fit <- HoltWinters(WWWusage, gamma=FALSE)
plot(forecast(fit))
```

---

`forecast.lm`


---

*Forecast a linear model with possible time series components*


---

**Description**

`forecast.lm` is used to predict linear models, especially those involving trend and seasonality components.

**Usage**

```
## S3 method for class 'lm'
forecast(object, newdata, level=c(80,95), fan=FALSE, h=10, ...)
```



**Arguments**

<code>object</code>	Object of class "lm", usually the result of a call to <code>lm</code> or <code>tslm</code> .
<code>newdata</code>	An optional data frame in which to look for variables with which to predict. If omitted, it is assumed that the only variables are trend and season, and <code>h</code> forecasts are produced.
<code>level</code>	Confidence level for prediction intervals.
<code>fan</code>	If TRUE, level is set to seq(50,99,by=1). This is suitable for fan plots.
<code>h</code>	Number of periods for forecasting. Ignored if <code>newdata</code> present.
<code>...</code>	Other arguments passed to <code>predict.lm()</code> .

**Details**

`forecast.lm` is largely a wrapper for `predict.lm()` except that it allows variables "trend" and "season" which are created on the fly from the time series characteristics of the data. Also, the output is reformatted into a `forecast` object.

**Value**

An object of class "forecast".

The function `summary` is used to obtain and print a summary of the results, while the function `plot` produces a plot of the forecasts and prediction intervals.

The generic accessor functions `fitted.values` and `residuals` extract useful features of the value returned by `forecast.lm`.

An object of class "forecast" is a list containing at least the following elements:

<code>model</code>	A list containing information about the fitted model
<code>method</code>	The name of the forecasting method as a character string
<code>mean</code>	Point forecasts as a time series
<code>lower</code>	Lower limits for prediction intervals
<code>upper</code>	Upper limits for prediction intervals
<code>level</code>	The confidence values associated with the prediction intervals
<code>x</code>	The historical data for the response variable.
<code>residuals</code>	Residuals from the fitted model. That is <code>x</code> minus fitted values.
<code>fitted</code>	Fitted values

**Author(s)**

Rob J Hyndman

**See Also**

`tslm`, `lm`.

## Examples

```
y <- ts(rnorm(120,0,3) + 1:120 + 20*sin(2*pi*(1:120)/12), frequency=12)
fit <- tslm(y ~ trend + season)
plot(forecast(fit, h=20))
```

---

forecast.stl	<i>Forecasting using stl objects</i>
--------------	--------------------------------------

---

## Description

Returns forecasts obtained by either ETS or ARIMA models applied to the seasonally adjusted data from an STL decomposition.

## Usage

```
## S3 method for class 'stl'
forecast(object, method=c("ets", "arima"),
         h = frequency(object$time.series)*2, level = c(80, 95), fan = FALSE, ...)
stlf(x, h=frequency(x)*2, s.window=7, method=c("ets", "arima"), level = c(80, 95), f
```

## Arguments

object	An object of class "stl". Usually the result of a call to <a href="#">stl</a> .
x	A univariate numeric time series of class "ts"
s.window	Either the character string "periodic" (default) or the span (in lags) of the loess window for seasonal extraction.
method	Method to use for forecasting the seasonally adjusted series.
h	Number of periods for forecasting.
level	Confidence level for prediction intervals.
fan	If TRUE, level is set to seq(50,99,by=1). This is suitable for fan plots.
...	Other arguments passed to <a href="#">ets()</a> or <a href="#">auto.arima()</a> .

## Details

`forecast.stl` seasonally adjusts the data from an STL decomposition, then uses either ETS or ARIMA models to forecast the result. The seasonal component from the last year of data is added back in to the forecasts. Note that the prediction intervals ignore the uncertainty associated with the seasonal component.

`stlf` takes a `ts` argument and applies a stl decomposition before calling `forecast.stl`.

**Value**

An object of class "forecast".

The function `summary` is used to obtain and print a summary of the results, while the function `plot` produces a plot of the forecasts and prediction intervals.

The generic accessor functions `fitted.values` and `residuals` extract useful features of the value returned by `forecast.stl`.

An object of class "forecast" is a list containing at least the following elements:

<code>model</code>	A list containing information about the fitted model
<code>method</code>	The name of the forecasting method as a character string
<code>mean</code>	Point forecasts as a time series
<code>lower</code>	Lower limits for prediction intervals
<code>upper</code>	Upper limits for prediction intervals
<code>level</code>	The confidence values associated with the prediction intervals
<code>x</code>	The original time series (either <code>object</code> itself or the time series used to create the model stored as <code>object</code> ).
<code>residuals</code>	Residuals from the fitted model. That is <code>x</code> minus fitted values.
<code>fitted</code>	Fitted values (one-step forecasts)

**Author(s)**

Rob J Hyndman

**See Also**

`forecast.ets`, `forecast.Arima`.

**Examples**

```
fit <- stl(USAccDeaths,s.window="periodic")
plot(forecast(fit))

plot(snaive(wineind))
```

---

`forecast.StructTS` *Forecasting using Structural Time Series models*

---

**Description**

Returns forecasts and other information for univariate structural time series models.

**Usage**

```
## S3 method for class 'StructTS'
forecast(object, h=ifelse(object$call$type=="BSM",2*object$xtsp[3],10),
         level=c(80,95), fan=FALSE, ...)
```

**Arguments**

<code>object</code>	An object of class "StructTS". Usually the result of a call to <a href="#">StructTS</a> .
<code>h</code>	Number of periods for forecasting
<code>level</code>	Confidence level for prediction intervals.
<code>fan</code>	If TRUE, level is set to seq(50,99,by=1). This is suitable for fan plots.
<code>...</code>	Other arguments.

**Details**

This function calls `predict.StructTS` and constructs an object of class "forecast" from the results.

**Value**

An object of class "forecast".

The function `summary` is used to obtain and print a summary of the results, while the function `plot` produces a plot of the forecasts and prediction intervals.

The generic accessor functions `fitted.values` and `residuals` extract useful features of the value returned by `forecast.StructTS`.

An object of class "forecast" is a list containing at least the following elements:

<code>model</code>	A list containing information about the fitted model
<code>method</code>	The name of the forecasting method as a character string
<code>mean</code>	Point forecasts as a time series
<code>lower</code>	Lower limits for prediction intervals
<code>upper</code>	Upper limits for prediction intervals
<code>level</code>	The confidence values associated with the prediction intervals
<code>x</code>	The original time series (either <code>object</code> itself or the time series used to create the model stored as <code>object</code> ).
<code>residuals</code>	Residuals from the fitted model. That is <code>x</code> minus fitted values.
<code>fitted</code>	Fitted values (one-step forecasts)

**Author(s)**

Rob J Hyndman

**See Also**

[StructTS](#).

**Examples**

```
fit <- StructTS(WWWusage, "level")
plot(forecast(fit))
```

---

`gas`*Australian monthly gas production*

---

**Description**

Australian monthly gas production: 1956–1995.

**Usage**`gas`**Format**

Time series data

**Source**

Australian Bureau of Statistics.

**Examples**

```
plot(gas)
seasonplot(gas)
tsdisplay(gas)
```

---

`gold`*Daily morning gold prices*

---

**Description**

Daily morning gold prices in US dollars. 1 January 1985 – 31 March 1989.

**Usage**`data(gold)`**Format**

Time series data

**Source**

Time Series Data Library. <http://robjhyndman.com/TSDL/>

**Examples**

```
tsdisplay(gold)
```

---

`logLik.ets`*Log-Likelihood of an ets object*

---

**Description**

Returns the log-likelihood of the ets model represented by `object` evaluated at the estimated parameters.

**Usage**

```
## S3 method for class 'ets'  
logLik(object, ...)
```

**Arguments**

<code>object</code>	an object of class <code>ets</code> , representing an exponential smoothing state space model.
<code>...</code>	some methods for this generic require additional arguments. None are used in this method.

**Value**

the log-likelihood of the model represented by `object` evaluated at the estimated parameters.

**Author(s)**

Rob J Hyndman

**References**

Hyndman, R.J., Koehler, A.B., Ord, J.K., and Snyder, R.D. (2008) *Forecasting with exponential smoothing: the state space approach*, Springer-Verlag. <http://robjhyndman.com/expsmooth>.

**See Also**

[ets](#)

**Examples**

```
fit <- ets(USAccDeaths)  
logLik(fit)
```

---

ma	<i>Moving-average smoothing</i>
----	---------------------------------

---

**Description**

Computes a simple moving average smoother.

**Usage**

```
ma(x, order, centre=TRUE)
```

**Arguments**

x	Univariate time series
order	Order of moving average smoother
centre	If TRUE, then the moving average is centred.

**Value**

Numerical time series object containing the smoothed values.

**Author(s)**

Rob J Hyndman

**See Also**

[ksmooth](#), [decompose](#)

**Examples**

```
plot(wineind)
sm <- ma(wineind, order=12)
lines(sm, col="red")
```

---

meanf	<i>Mean Forecast</i>
-------	----------------------

---

**Description**

Returns forecasts and prediction intervals for an iid model applied to x.

**Usage**

```
meanf(x, h=10, level=c(80, 95), fan=FALSE)
```

**Arguments**

<code>x</code>	a numeric vector or time series
<code>h</code>	Number of periods for forecasting
<code>level</code>	Confidence levels for prediction intervals.
<code>fan</code>	If TRUE, level is set to seq(50,99,by=1). This is suitable for fan plots.

**Details**

The iid model is

$$Y_t = \mu + Z_t$$

where  $Z_t$  is a normal iid error. Forecasts are given by

$$Y_n(h) = \mu$$

where  $\mu$  is estimated by the sample mean.

**Value**

An object of class "forecast".

The function `summary` is used to obtain and print a summary of the results, while the function `plot` produces a plot of the forecasts and prediction intervals.

The generic accessor functions `fitted.values` and `residuals` extract useful features of the value returned by `meanf`.

An object of class "forecast" is a list containing at least the following elements:

<code>model</code>	A list containing information about the fitted model
<code>method</code>	The name of the forecasting method as a character string
<code>mean</code>	Point forecasts as a time series
<code>lower</code>	Lower limits for prediction intervals
<code>upper</code>	Upper limits for prediction intervals
<code>level</code>	The confidence values associated with the prediction intervals
<code>x</code>	The original time series (either <code>object</code> itself or the time series used to create the model stored as <code>object</code> ).
<code>residuals</code>	Residuals from the fitted model. That is <code>x</code> minus fitted values.
<code>fitted</code>	Fitted values (one-step forecasts)

**Author(s)**

Rob J Hyndman

**See Also**

[rwf](#)

**Examples**

```
nile.fcast <- meanf(Nile, h=10)
plot(nile.fcast)
```



---

monthdays	<i>Number of days in each season</i>
-----------	--------------------------------------

---

**Description**

Returns number of days in each month or quarter of the observed time period.

**Usage**

```
monthdays(x)
```

**Arguments**

`x`                      time series

**Details**

Useful for month length adjustments

**Value**

Time series

**Author(s)**

Rob J Hyndman

**Examples**

```
par(mfrow=c(2,1))
plot(ldeaths,xlab="Year",ylab="pounds",
     main="Monthly deaths from lung disease (UK)")
ldeaths.adj <- ldeaths/monthdays(ldeaths)*365.25/12
plot(ldeaths.adj,xlab="Year",ylab="pounds",
     main="Adjusted monthly deaths from lung disease (UK)")
```

---

na.interp	<i>Interpolate missing values in a time series</i>
-----------	--

---

**Description**

Uses linear interpolation to replace missing values.

**Usage**

```
na.interp(x)
```

**Arguments**

`x` time series

**Details**

A more general and flexible approach is available using `na.approx` in the `zoo` package.

**Value**

Time series

**Author(s)**

Rob J Hyndman

**Examples**

```
data(gold)
plot(na.interp(gold))
```

---

naive

*Naive forecasts*

---

**Description**

`naive()` returns forecasts and prediction intervals for an ARIMA(0,1,0) random walk model applied to `x`. `snaive()` returns forecasts and prediction intervals from an ARIMA(0,0,0)(0,1,0)<sub>m</sub> model where `m` is the seasonal period.

**Usage**

```
naive(x, h=10, level=c(80,95), fan=FALSE)
snaive(x, h=2*frequency(x), level=c(80,95), fan=FALSE)
```

**Arguments**

`x` a numeric vector or time series

`h` Number of periods for forecasting

`level` Confidence levels for prediction intervals.

`fan` If TRUE, level is set to `seq(50,99,by=1)`. This is suitable for fan plots.

**Details**

These functions are simply convenient wrappers to [Arima](#) with the appropriate arguments to return naive and seasonal naive forecasts.

**Value**

An object of class "forecast".

The function `summary` is used to obtain and print a summary of the results, while the function `plot` produces a plot of the forecasts and prediction intervals.

The generic accessor functions `fitted.values` and `residuals` extract useful features of the value returned by `naive` or `snaive`.

An object of class "forecast" is a list containing at least the following elements:

<code>model</code>	A list containing information about the fitted model
<code>method</code>	The name of the forecasting method as a character string
<code>mean</code>	Point forecasts as a time series
<code>lower</code>	Lower limits for prediction intervals
<code>upper</code>	Upper limits for prediction intervals
<code>level</code>	The confidence values associated with the prediction intervals
<code>x</code>	The original time series (either <code>object</code> itself or the time series used to create the model stored as <code>object</code> ).
<code>residuals</code>	Residuals from the fitted model. That is <code>x</code> minus fitted values.
<code>fitted</code>	Fitted values (one-step forecasts)

**Author(s)**

Rob J Hyndman

**See Also**

[Arima](#), [rwf](#)

**Examples**

```
plot(naive(gold,h=50),include=200)
plot(snaive(wineind))
```

---

ndiffs

*Number of differences*


---

**Description**

Uses a unit root test to determine the number of differences required for time series `x`. If `test="kpss"`, the KPSS test is used with the null hypothesis that `x` has a stationary root against a unit-root alternative. Then the test returns the least number of differences required to pass the test at the level `alpha`. If `test="adf"`, the Augmented Dickey-Fuller test is used and if `test="pp"` the Phillips-Perron test is used. In both of these cases, the null hypothesis is that `x` has a unit root against a stationary root alternative. Then the test returns the least number of differences required to fail the test at the level `alpha`.

**Usage**

```
ndiffs(x, alpha=0.05, test=c("kpss", "adf", "pp"))
```

**Arguments**

x	A univariate time series
alpha	Level of the test
test	Type of unit root test to use

**Value**

An integer.

**Author(s)**

Rob J Hyndman

**Examples**

```
ndiffs(WWWusage)
```

---

plot.ets

*Plot components from ETS model*

---

**Description**

Produces a plot of the level, slope and seasonal components from an ETS model.

**Usage**

```
## S3 method for class 'ets'
plot(x, ...)
```

**Arguments**

x	Object of class “ets”.
...	Other plotting parameters passed to <a href="#">par</a> .

**Value**

None. Function produces a plot

**Author(s)**

Rob J Hyndman

**See Also**[ets](#)**Examples**

```
fit <- ets(USAccDeaths)
plot(fit)
plot(fit, plot.type="single", ylab="", col=1:3)
```

---

plot.forecast	<i>Forecast plot</i>
---------------	----------------------

---

**Description**

Plots a time series with forecasts and prediction intervals.

**Usage**

```
## S3 method for class 'forecast'
plot(x, include, plot.conf = TRUE, shaded = TRUE, shadebars=(length(x$mean)<5),
      shadecols=NULL, lambda = NULL, col = 1, fcol = 4, pi.col=1, pi.lty
      ylim = NULL, main = NULL, ylab = "", xlab = "", ...)

## S3 method for class 'splineforecast'
plot(x, fitcol=2, ...)
```

**Arguments**

<code>x</code>	Forecast object produced by <a href="#">forecast</a> .
<code>include</code>	number of values from time series to include in plot
<code>plot.conf</code>	Logical flag indicating whether to plot prediction intervals.
<code>shaded</code>	Logical flag indicating whether prediction intervals should be shaded (TRUE) or lines (FALSE)
<code>shadebars</code>	Logical flag indicating if prediction intervals should be plotted as shaded bars (if TRUE) or a shaded polygon (if FALSE). Ignored if <code>shaded=FALSE</code> . Bars are plotted by default if there are fewer than five forecast horizons.
<code>shadecols</code>	Colors for shaded prediction intervals
<code>lambda</code>	Box-Cox parameter. If present, function backtransforms data and forecasts using <a href="#">InvBoxCox</a> .
<code>col</code>	the colour for the data line.
<code>fcol</code>	the colour for the forecast line.
<code>pi.col</code>	If <code>shade=FALSE</code> and <code>plot.conf=TRUE</code> , the prediction intervals are plotted in this colour.

<code>pi.lty</code>	If <code>shade=FALSE</code> and <code>plot.conf=TRUE</code> , the prediction intervals are plotted using this line type.
<code>ylim</code>	Limits on y-axis
<code>main</code>	Main title
<code>ylab</code>	Y-axis label
<code>xlab</code>	X-axis label
<code>fitcol</code>	Line colour for fitted values.
<code>...</code>	additional arguments to <code>plot</code> .

**Value**

None.

**Author(s)**

Rob J Hyndman

**References**

Makridakis, Wheelwright and Hyndman (1998) *Forecasting: methods and applications*, Wiley: New York. <http://robjhyndman.com/forecasting>.

**See Also**

`plot.ts`

**Examples**

```
deaths.fit <- hw(USAccDeaths,h=48)
plot(deaths.fit)
```

---

`rwf`

*Random Walk Forecast*

---

**Description**

Returns forecasts and prediction intervals for a random walk with drift model applied to `x`.

**Usage**

```
rwf(x, h=10, drift=FALSE, level=c(80,95), fan=FALSE)
```

**Arguments**

<code>x</code>	a numeric vector or time series
<code>h</code>	Number of periods for forecasting
<code>drift</code>	Logical flag. If TRUE, fits a random walk with drift model.
<code>level</code>	Confidence levels for prediction intervals.
<code>fan</code>	If TRUE, level is set to seq(50,99,by=1). This is suitable for fan plots.

**Details**

The random walk with drift model is

$$Y_t = c + Y_{t-1} + Z_t$$

where  $Z_t$  is a normal iid error. Forecasts are given by

$$Y_n(h) = ch + Y_n$$

. If there is no drift, the drift parameter  $c=0$ . Forecast standard errors allow for uncertainty in estimating the drift parameter.

**Value**

An object of class "forecast".

The function `summary` is used to obtain and print a summary of the results, while the function `plot` produces a plot of the forecasts and prediction intervals.

The generic accessor functions `fitted.values` and `residuals` extract useful features of the value returned by `rwf`.

An object of class "forecast" is a list containing at least the following elements:

<code>model</code>	A list containing information about the fitted model
<code>method</code>	The name of the forecasting method as a character string
<code>mean</code>	Point forecasts as a time series
<code>lower</code>	Lower limits for prediction intervals
<code>upper</code>	Upper limits for prediction intervals
<code>level</code>	The confidence values associated with the prediction intervals
<code>x</code>	The original time series (either <code>object</code> itself or the time series used to create the model stored as <code>object</code> ).
<code>residuals</code>	Residuals from the fitted model. That is <code>x</code> minus fitted values.
<code>fitted</code>	Fitted values (one-step forecasts)

**Author(s)**

Rob J Hyndman

**See Also**

[arima](#), [meanf](#)

**Examples**

```
gold.fcast <- rwf(gold[1:60],h=50)
plot(gold.fcast)
```

---

seasadj

*Seasonal adjustment*

---

**Description**

Returns seasonally adjusted data constructed by removing the seasonal component.

**Usage**

```
seasadj(object)
```

**Arguments**

object            Object created by [decompose](#) or [stl](#).

**Value**

Univariate time series.

**Author(s)**

Rob J Hyndman

**See Also**

[stl](#), [decompose](#)

**Examples**

```
plot(AirPassengers)
lines(seasadj(decompose(AirPassengers,"multiplicative")),col=4)
```



---

seasonaldummy	<i>Seasonal dummy variables</i>
---------------	---------------------------------

---

## Description

`seasonaldummy` and `seasonaldummyf` return matrices of dummy variables suitable for use in `arima`, `lm` or `tslm`. The last season is omitted and used as the control.

`fourier` and `fourierf` return matrices containing terms from a Fourier series, up to order  $K$ , suitable for use in `arima`, `lm` or `tslm`.

## Usage

```
seasonaldummy(x)
seasonaldummyf(x, h)
fourier(x, K)
fourierf(x, K, h)
```

## Arguments

<code>x</code>	Seasonal time series
<code>h</code>	Number of periods ahead to forecast
<code>K</code>	Maximum order of Fourier terms

## Value

Numerical matrix with number of rows equal to the `length(x)` and number of columns equal to `frequency(x)-1` (for `seasonaldummy` and `seasonaldummyf` or  $2*K$  (for `fourier` or `fourierf`).

## Author(s)

Rob J Hyndman

## Examples

```
plot(ldeaths)

# Using seasonal dummy variables
month <- seasonaldummy(ldeaths)
deaths.lm <- tslm(ldeaths ~ month)
tsdisplay(residuals(deaths.lm))
ldeaths.fcast <- forecast(deaths.lm, data.frame(month=I(seasonaldummyf(ldeaths,36))))
plot(ldeaths.fcast)

# A simpler approach to seasonal dummy variables
deaths.lm <- tslm(ldeaths ~ season)
ldeaths.fcast <- forecast(deaths.lm, h=36)
plot(ldeaths.fcast)
```

```
# Using Fourier series
X <- fourier(ldeaths,3)
deaths.lm <- tslm(ldeaths ~ X)
ldeaths.fcast <- forecast(deaths.lm, data.frame(X=I(fourierf(ldeaths,3,36))))
plot(ldeaths.fcast)
```

---

seasonplot

*Seasonal plot*


---

### Description

Plots a seasonal plot as described in Makridakis, Wheelwright and Hyndman (1998, chapter 2).

### Usage

```
seasonplot(x, s, season.labels = NULL, year.labels = FALSE,
           year.labels.left = FALSE, type = "o", main, ylab = "",
           xlab = NULL, col = 1, ...)
```

### Arguments

<code>x</code>	a numeric vector or time series.
<code>s</code>	seasonal frequency of <code>x</code>
<code>season.labels</code>	Labels for each season in the "year"
<code>year.labels</code>	Logical flag indicating whether labels for each year of data should be plotted on the right.
<code>year.labels.left</code>	Logical flag indicating whether labels for each year of data should be plotted on the left.
<code>type</code>	plot type (as for <a href="#">plot</a> )
<code>main</code>	Main title.
<code>ylab</code>	Y-axis label
<code>xlab</code>	X-axis label
<code>col</code>	Colour
<code>...</code>	additional arguments to <a href="#">plot</a> .

### Value

None.

### Author(s)

Rob J Hyndman

## References

Makridakis, Wheelwright and Hyndman (1998) *Forecasting: methods and applications*, Wiley: New York. <http://robjhyndman.com/forecasting/>

## See Also

`monthplot`

## Examples

```
seasonplot(AirPassengers)
```

---

ses

*Exponential smoothing forecasts*

---

## Description

Returns forecasts and other information for exponential smoothing forecasts applied to x.

## Usage

```
ses(x, h=10, level=c(80,95), fan=FALSE, ...)
```

```
holt(x, h=10, damped=FALSE, level=c(80,95), fan=FALSE, ...)
```

```
hw(x, h=2*frequency(x), seasonal="additive", damped=FALSE, level=c(80,95), fan=FALSE, ...)
```

## Arguments

x	a numeric vector or time series
h	Number of periods for forecasting.
damped	If TRUE, use a damped trend.
seasonal	Type of seasonality in hw model. "additive" or "multiplicative"
level	Confidence level for prediction intervals.
fan	If TRUE, level is set to seq(50,99,by=1). This is suitable for fan plots.
...	Other arguments passed to <code>ets</code> .

## Details

ses, holt and hw are simply convenient wrapper functions for `forecast(ets(...))`.

**Value**

An object of class "forecast".

The function `summary` is used to obtain and print a summary of the results, while the function `plot` produces a plot of the forecasts and prediction intervals.

The generic accessor functions `fitted.values` and `residuals` extract useful features of the value returned by `ets` and associated functions.

An object of class "forecast" is a list containing at least the following elements:

<code>model</code>	A list containing information about the fitted model
<code>method</code>	The name of the forecasting method as a character string
<code>mean</code>	Point forecasts as a time series
<code>lower</code>	Lower limits for prediction intervals
<code>upper</code>	Upper limits for prediction intervals
<code>level</code>	The confidence values associated with the prediction intervals
<code>x</code>	The original time series (either <code>object</code> itself or the time series used to create the model stored as <code>object</code> ).
<code>residuals</code>	Residuals from the fitted model. That is <code>x</code> minus fitted values.
<code>fitted</code>	Fitted values (one-step forecasts)

**Author(s)**

Rob J Hyndman

**References**

Hyndman, R.J., Koehler, A.B., Snyder, R.D., Grose, S. (2002) "A state space framework for automatic forecasting using exponential smoothing methods", *International J. Forecasting*, **18**(3), 439–454.

Hyndman, R.J., Akram, Md., and Archibald, B. (2008) "The admissible parameter space for exponential smoothing models". *Annals of Statistical Mathematics*, **60**(2), 407–426.

**See Also**

[`ets`](#), [`HoltWinters`](#), [`rwf`](#), [`arima`](#).

**Examples**

```
fcast <- holt(airmiles)
plot(fcast)
deaths.fcast <- hw(USAccDeaths,h=48)
plot(deaths.fcast)
```

simulate.ets

*Simulation from a time series model***Description**

Returns a time series based on the model object `object`.

**Usage**

```
## S3 method for class 'ets'
simulate(object, nsim=length(object$x), seed=NULL, future=TRUE, bootstrap=FALSE, ...)
## S3 method for class 'ar'
simulate(object, nsim=object$n.used, seed=NULL, future=TRUE, bootstrap=FALSE, ...)
## S3 method for class 'Arima'
simulate(object, nsim=length(object$x), seed=NULL, xreg=NULL, future=TRUE, bootstrap=FALSE, ...)
## S3 method for class 'fracdiff'
simulate(object, nsim=object$n, seed=NULL, future=TRUE, bootstrap=FALSE, ...)
```

**Arguments**

<code>object</code>	An object of class "ets", "Arima" or "ar".
<code>nsim</code>	Number of periods for the simulated series
<code>seed</code>	Either NULL or an integer that will be used in a call to <code>set.seed</code> before simulating the time series. The default, NULL will not change the random generator state.
<code>future</code>	Produce sample paths that are future to and conditional on the data in <code>object</code> .
<code>bootstrap</code>	If TRUE, simulation uses resampled errors rather than normally distributed errors.
<code>xreg</code>	New values of <code>xreg</code> to be used for forecasting. Must have <code>nsim</code> rows.
<code>...</code>	Other arguments.

**Value**

An object of class "ts".

**Author(s)**

Rob J Hyndman

**See Also**

`ets`, `Arima`, `auto.arima`, `ar`, `arfima`.

**Examples**

```
fit <- ets(USAccDeaths)
plot(USAccDeaths, xlim=c(1973, 1982))
lines(simulate(fit, 36), col="red")
```

sindexf

*Forecast seasonal index***Description**

Returns vector containing the seasonal index for  $h$  future periods. If the seasonal index is non-periodic, it uses the last values of the index.

**Usage**

```
sindexf(object, h)
```

**Arguments**

object	Output from <a href="#">decompose</a> or <a href="#">stl</a> .
h	Number of periods ahead to forecast

**Value**

Time series

**Author(s)**

Rob J Hyndman

**Examples**

```
uk.stl <- stl(UKDriverDeaths, "periodic")
uk.sa <- seasadj(uk.stl)
uk.fcast <- holt(uk.sa, 36)
seasf <- sindexf(uk.stl, 36)
uk.fcast$mean <- uk.fcast$mean + seasf
uk.fcast$lower <- uk.fcast$lower + cbind(seasf, seasf)
uk.fcast$upper <- uk.fcast$upper + cbind(seasf, seasf)
uk.fcast$x <- UKDriverDeaths
plot(uk.fcast, main="Forecasts from Holt's method with seasonal adjustment")
```

splinef

*Cubic Spline Forecast***Description**

Returns local linear forecasts and prediction intervals using cubic smoothing splines.

**Usage**

```
splinef(x, h=10, level=c(80,95), fan=FALSE)
```

**Arguments**

x	a numeric vector or time series
h	Number of periods for forecasting
level	Confidence level for prediction intervals.
fan	If TRUE, level is set to seq(50,99,by=1). This is suitable for fan plots.

**Details**

The cubic smoothing spline model is equivalent to an ARIMA(0,2,2) model but with a restricted parameter space. The advantage of the spline model over the full ARIMA model is that it provides a smooth historical trend as well as a linear forecast function. Hyndman, King, Pitrun, and Billah (2002) show that the forecast performance of the method is hardly affected by the restricted parameter space.

**Value**

An object of class "forecast".

The function `summary` is used to obtain and print a summary of the results, while the function `plot` produces a plot of the forecasts and prediction intervals.

The generic accessor functions `fitted.values` and `residuals` extract useful features of the value returned by `meanf`.

An object of class "forecast" is a list containing at least the following elements:

model	A list containing information about the fitted model
method	The name of the forecasting method as a character string
mean	Point forecasts as a time series
lower	Lower limits for prediction intervals
upper	Upper limits for prediction intervals
level	The confidence values associated with the prediction intervals
x	The original time series (either <code>object</code> itself or the time series used to create the model stored as <code>object</code> ).
residuals	Residuals from the fitted model. That is <code>x</code> minus fitted values.
fitted	Fitted values (one-step forecasts)

**Author(s)**

Rob J Hyndman

**References**

Hyndman, King, Pitrun and Billah (2005) Local linear forecasts using cubic smoothing splines. *Australian and New Zealand Journal of Statistics*, **47**(1), 87-99. <http://robjhyndman.com/papers/splinefcast.htm>.

**See Also**

[smooth.spline](#), [arima](#), [holt](#).

**Examples**

```
fcast <- splinef(uspop,h=5)
plot(fcast)
summary(fcast)
```

---

thetaf

*Theta method forecast*


---

**Description**

Returns forecasts and prediction intervals for a theta method forecast.

**Usage**

```
thetaf(x, h=10, level=c(80,95), fan=FALSE)
```

**Arguments**

x	a numeric vector or time series
h	Number of periods for forecasting
level	Confidence levels for prediction intervals.
fan	If TRUE, level is set to seq(50,99,by=1). This is suitable for fan plots.

**Details**

The theta method of Assimakopoulos and Nikolopoulos (2000) is equivalent to simple exponential smoothing with drift. This is demonstrated in Hyndman and Billah (2003). Prediction intervals are computed using the underlying state space model.



**Value**

An object of class "forecast".

The function `summary` is used to obtain and print a summary of the results, while the function `plot` produces a plot of the forecasts and prediction intervals.

The generic accessor functions `fitted.values` and `residuals` extract useful features of the value returned by `rwf`.

An object of class "forecast" is a list containing at least the following elements:

<code>model</code>	A list containing information about the fitted model
<code>method</code>	The name of the forecasting method as a character string
<code>mean</code>	Point forecasts as a time series
<code>lower</code>	Lower limits for prediction intervals
<code>upper</code>	Upper limits for prediction intervals
<code>level</code>	The confidence values associated with the prediction intervals
<code>x</code>	The original time series (either <code>object</code> itself or the time series used to create the model stored as <code>object</code> ).
<code>residuals</code>	Residuals from the fitted model. That is <code>x</code> minus fitted values.
<code>fitted</code>	Fitted values (one-step forecasts)

**Author(s)**

Rob J Hyndman

**References**

Assimakopoulos, V. and Nikolopoulos, K. (2000). The theta model: a decomposition approach to forecasting. *International Journal of Forecasting* **16**, 521-530.

Hyndman, R.J., and Billah, B. (2003) Unmasking the Theta method. *International J. Forecasting*, **19**, 287-290.

**See Also**

`arima`, `meanf`, `rwf`, `ses`

**Examples**

```
nile.fcast <- thetaf(Nile)
plot(nile.fcast)
```

tsdisplay

*Time series display***Description**

Plots a time series along with its acf and either its pacf, lagged scatterplot or spectrum.

**Usage**

```
tsdisplay(x, plot.type = "partial", points = TRUE, ci.type =
          "white", lag.max = round(10 * log10(length(x))),
          na.action = na.interp, main = NULL, ylab = "", xlab =
          "", pch = 1, cex = 0.5, ...)
```

**Arguments**

<code>x</code>	a numeric vector or time series.
<code>plot.type</code>	type of plot to include in lower right corner. Possible values are "partial", "scatter" or "spectrum".
<code>points</code>	logical flag indicating whether to show the individual points or not in the time plot.
<code>ci.type</code>	type of confidence limits for ACF. Possible values are as for <a href="#">acf</a> .
<code>lag.max</code>	the maximum lag to plot for the acf and pacf.
<code>na.action</code>	how to handle missing values. Default is to use linear interpolation.
<code>main</code>	Main title.
<code>ylab</code>	Y-axis label
<code>xlab</code>	X-axis label
<code>pch</code>	Plotting character
<code>cex</code>	Character size
<code>...</code>	additional arguments to <a href="#">acf</a> .

**Value**

None.

**Author(s)**

Rob J Hyndman

**References**

Makridakis, Wheelwright and Hyndman (1998) *Forecasting: methods and applications*, Wiley: New York. <http://robjhyndman.com/forecasting/>

**See Also**

[plot.ts](#), [acf](#)

**Examples**

```
tsdisplay(diff(WWWusage))
```

---

tslm

---

*Fit a linear model with time series components*


---

**Description**

`tslm` is used to fit linear models to time series including trend and seasonality components.

**Usage**

```
tslm(formula, data, ...)
```

**Arguments**

<code>formula</code>	an object of class "formula" (or one that can be coerced to that class): a symbolic description of the model to be fitted.
<code>data</code>	an optional data frame, list or environment (or object coercible by <code>as.data.frame</code> to a data frame) containing the variables in the model. If not found in data, the variables are taken from <code>environment(formula)</code> , typically the environment from which <code>lm</code> is called.
<code>...</code>	Other arguments passed to <code>lm()</code> .

**Details**

`tslm` is largely a wrapper for `lm()` except that it allows variables "trend" and "season" which are created on the fly from the time series characteristics of the data. The variable "trend" is a simple time trend and "season" is a factor indicating the season (e.g., the month or the quarter depending on the frequency of the data).

**Value**

Returns an object of class "lm".

**Author(s)**

Rob J Hyndman

**See Also**

[forecast.lm](#), [lm](#).

**Examples**

```
y <- ts(rnorm(120,0,3) + 1:120 + 20*sin(2*pi*(1:120)/12), frequency=12)
fit <- tslm(y ~ trend + season)
plot(forecast(fit, h=20))
```

---

wineind

---

*Australian total wine sales*


---

**Description**

Australian total wine sales by wine makers in bottles <= 1 litre. Jan 1980 – Aug 1994.

**Usage**

```
wineind
```

**Format**

Time series data

**Source**

Time Series Data Library. <http://robjhyndman.com/TSDL/>

**Examples**

```
tsdisplay(wineind)
```

---

woolryrq

---

*Quarterly production of woollen yarn in Australia*


---

**Description**

Quarterly production of woollen yarn in Australia: tonnes. Mar 1965 – Sep 1994.

**Usage**

```
woolryrq
```

**Format**

Time series data

**Source**

Time Series Data Library. <http://robjhyndman.com/TSDL/>

*woolyrnq*

53

### **Examples**

```
tsdisplay(woolyrnq)
```

# Index

## \*Topic **datasets**

gas, 29  
gold, 29  
wineind, 52  
woolryrnq, 52

## \*Topic **hplot**

plot.ets, 36

## \*Topic **htest**

dm.test, 15

## \*Topic **models**

CV, 13

## \*Topic **stats**

forecast.lm, 24  
tslm, 51

## \*Topic **ts**

accuracy, 2  
Acf, 3  
arfima, 4  
Arima, 6  
arima.errors, 8  
auto.arima, 9  
BoxCox, 10  
croston, 11  
decompose, 14  
dm.test, 15  
ets, 16  
fitted.Arima, 18  
forecast, 19  
forecast.Arima, 20  
forecast.ets, 22  
forecast.HoltWinters, 23  
forecast.stl, 26  
forecast.StructTS, 27  
logLik.ets, 30  
ma, 31  
meanf, 31  
monthdays, 33  
na.interp, 33  
naive, 34

ndiffs, 35  
plot.forecast, 37  
rwf, 38  
seasadj, 40  
seasonaldummy, 41  
seasonplot, 42  
ses, 43  
simulate.ets, 45  
sindexf, 46  
splinef, 47  
thetaf, 48  
tsdisplay, 50

accuracy, 2  
Acf, 3  
acf, 3, 4, 50, 51  
AIC, 13  
ar, 20, 21, 45  
arfima, 4, 20, 21, 45  
Arima, 6, 10, 21, 34, 35, 45  
arima, 5–8, 10, 18–21, 40, 41, 44, 48, 49  
arima.errors, 8  
auto.arima, 5, 9, 20, 21, 26, 45

best.arima(auto.arima), 9  
BoxCox, 10

croston, 11, 20  
CV, 13

decompose, 14, 14, 15, 31, 40, 46  
dm.test, 15

ets, 16, 19, 22, 23, 26, 30, 37, 44, 45

fitted.Arima, 18  
forecast, 19, 37  
forecast.ar(forecast.Arima), 20  
forecast.Arima, 18, 19, 20, 20, 27  
forecast.ets, 19, 20, 22, 27  
forecast.fracdiff, 5, 21

`forecast.fracdiff`  
    (`forecast.Arima`), 20  
`forecast.HoltWinters`, 20, 23  
`forecast.lm`, 24, 51  
`forecast.stl`, 26  
`forecast.StructTS`, 20, 27  
`forecast.ts`, 19  
`fourier(seasonaldummy)`, 41  
`fourierf(seasonaldummy)`, 41  
`fracdiff`, 5, 20  
  
`gas`, 29  
`gold`, 29  
  
`holt`, 20, 48  
`holt(ses)`, 43  
`HoltWinters`, 18, 23, 24, 44  
`hw`, 20  
`hw(ses)`, 43  
  
`InvBoxCox`, 3, 37  
`InvBoxCox(BoxCox)`, 10  
  
`ksmooth`, 31  
  
`lm`, 13, 25, 41, 51  
`logLik.ets`, 30  
  
`ma`, 31  
`meanf`, 20, 31, 40, 49  
`monthdays`, 33  
`monthplot`, 43  
  
`na.interp`, 33  
`naive`, 34  
`ndiffs`, 10, 35  
  
`Pacf(Acf)`, 3  
`par`, 36  
`plot`, 38, 42  
`plot.decomposed.ts(decompose)`, 14  
`plot.ets`, 36  
`plot.forecast`, 37  
`plot.splineforecast`  
    (`plot.forecast`), 37  
`plot.ts`, 38, 51  
`predict.ar`, 21  
`predict.Arima`, 21  
`predict.HoltWinters`, 23, 24  
`predict.lm`, 25  
  
`print.forecast(forecast)`, 19  
  
`residuals`, 8  
`rwf`, 18, 20, 32, 35, 38, 44, 49  
  
`seasadj`, 40  
`seasonaldummy`, 41  
`seasonaldummyf(seasonaldummy)`, 41  
`seasonplot`, 42  
`ses`, 12, 20, 43, 49  
`set.seed`, 45  
`simulate.ar(simulate.ets)`, 45  
`simulate.Arima(simulate.ets)`, 45  
`simulate.ets`, 45  
`simulate.fracdiff(simulate.ets)`,  
    45  
`sindexf`, 46  
`smooth.spline`, 48  
`snaive(naive)`, 34  
`splinef`, 20, 47  
`stl`, 26, 40, 46  
`stl(forecast.stl)`, 26  
`stlf(forecast.stl)`, 26  
`StructTS`, 28  
`summary.forecast(forecast)`, 19  
  
`thetaf`, 20, 48  
`tsdisplay`, 50  
`tslm`, 13, 25, 41, 51  
  
`wineind`, 52  
`woolyrnq`, 52